

10301/601 Homework 2 (Written)

Joong Ho Choi

TOTAL POINTS

27 / 30

QUESTION 1

1 Bonus 1 / 0

- ✓ + 1 pts Correct- this has been added separately and will appear on your grade summary
- 1 pts Incorrect

2.7 2.7 1 / 1

- ✓ - 0 pts Correct
- 1 pts Incorrect
- 1 pts Blank

QUESTION 2

Warm-Up 11 pts

2.1 2.1 1 / 1

- ✓ - 0 pts Correct: 1
- 1 pts Incorrect

2.2 2.2 1 / 1

- ✓ - 0 pts Correct
- 1 pts Incorrect
- 0.5 pts Should include number rounded to 4th decimal place

2.8 2.8 1 / 4

- ✓ + 1 pts first split is correct: C at root
- + 1 pts second split is correct: C=0 --> A
- + 1 pts third split is correct: A=0 --> B
- + 1 pts leaf nodes are correct:
 - C=1 --> Y=0
 - A=1 --> Y=1
 - B=1 --> Y=0
 - B=2 --> Y=1
- + 0 pts no nodes correct

2.3 2.3 1 / 1

- ✓ - 0 pts Correct
- 1 pts Wrong/Blank
- 0.5 pts No/Wrong rounding

QUESTION 3

Pseudocode 6 pts

2.4 2.4 1 / 1

- ✓ - 0 pts Correct
- 1 pts Incorrect
- 1 pts Blank

3.1 3.1 3 / 3

- ✓ - 0 pts Correct
- 1 pts wrong if statement
- 1 pts wrong return value
- 3 pts incorrect/blank

2.5 2.5 1 / 1

- ✓ - 0 pts Correct
- 1 pts Incorrect

3.2 3.2 3 / 3

- ✓ - 0 pts Correct
- 1 pts wrong casing
- 1 pts wrong recursive call
- 3 pts incorrect/blank

2.6 2.6 0 / 1

- 0 pts Correct
- ✓ - 1 pts Incorrect

QUESTION 4

Empirical 13 pts

4.1 4.1.1 1 / 1

- ✓ - 0 pts Correct

- 1 pts Blank/Incorrect

4.2 4.1.2 1 / 1

✓ - 0 pts Correct

- 1 pts Incorrect/Blank

4.3 4.1.3 1 / 1

✓ - 0 pts Correct

- 1 pts Incorrect/Blank

4.4 4.1.4 1 / 1

✓ - 0 pts Correct

- 1 pts Incorrect/Blank

4.5 4.2 3 / 3

✓ - 0 pts Correct trends for both test and train

- 3 pts Blank/Missing graph

- 3 pts Wrong trends for both test and train

- 1 pts Wrong graph types/plot on separate graphs

- 1.5 pts Too few points on the graph

4.6 4.3 2 / 2

✓ - 0 pts Correct: D

- 1 pts Correct parity with incorrect justification

- 2 pts Incorrect / Blank

4.7 4.4 2 / 2

✓ - 0 pts Correct: A

- 2 pts Incorrect / Blank

4.8 4.5 2 / 2

✓ - 0 pts Correct

- 1 pts Minor Error

- 2 pts Blank

QUESTION 5

5 Collaboration 0 / 0

✓ - 0 pts Correct

HOMEWORK 2: DECISION TREES

10-301/10-601 Introduction to Machine Learning (Spring 2022)

<http://www.cs.cmu.edu/~mgormley/courses/10601/>

OUT: Wednesday, January 26th

DUE: Friday, February 4th

TAs: Zachary, Sami, Yuxin, Junhui, Rita

Summary It's time to build your first end-to-end learning system! In this assignment, you will build a Decision Tree classifier and apply it to several binary classification problems. This assignment consists of several parts: In the Written component, you will work through some Information Theory basics in order to "learn" a Decision Tree on paper, and also work through some pseudocode that will help you algorithmically think through the programming assignment. Then in Programming component, you will implement Decision Tree learning, prediction, and evaluation. Using that implementation, you will answer the empirical questions found at the end of the Written component.

START HERE: Instructions

- **Collaboration Policy:** Please read the collaboration policy here: <http://www.cs.cmu.edu/~mgormley/courses/10601/syllabus.html>
- **Late Submission Policy:** See the late submission policy here: <http://www.cs.cmu.edu/~mgormley/courses/10601/syllabus.html>
- **Submitting your work:** You will use Gradescope to submit answers to all questions and code. Please follow instructions at the end of this PDF to correctly submit all your code to Gradescope.
 - **Written:** For written problems such as short answer, multiple choice, derivations, proofs, or plots, please use the provided template. Submissions can be handwritten onto the template, but should be labeled and clearly legible. If your writing is not legible, you will not be awarded marks. If your scanned submission misaligns the template, there will be a 5% penalty. Alternatively, submissions can be written in LaTeX. Each derivation/proof should be completed in the boxes provided. If you do not follow the template, your assignment may not be graded correctly by our AI assisted grader.
 - **Programming:** You will submit your code for programming questions on the homework to Gradescope (<https://gradescope.com>). After uploading your code, our grading scripts will autograde your assignment by running your program on a virtual machine (VM). When you are developing, check that the version number of the programming language environment (e.g. Python 3.9.6) and versions of permitted libraries (e.g. numpy 1.21.2 and scipy 1.7.1) match those used on Gradescope. You have 10 free Gradescope programming submissions. After 10 submissions, you will begin to lose points from your total programming score. We recommend debugging your implementation on your local machine (or the Linux servers) and making sure your code is running correctly first before submitting your code to Gradescope.

- **Materials:** The data that you will need in order to complete this assignment is posted along with the writeup and template on the course website.

Instructions for Specific Problem Types

For “Select One” questions, please fill in the appropriate bubble completely:

Select One: Who taught this course?

- Matt Gormley
- Marie Curie
- Noam Chomsky

If you need to change your answer, you may cross out the previous answer and bubble in the new answer:

Select One: Who taught this course?

- Matt Gormley
- Marie Curie
- Noam Chomsky

For “Select all that apply” questions, please fill in all appropriate squares completely:

Select all that apply: Which are scientists?

- Stephen Hawking
- Albert Einstein
- Isaac Newton
- I don't know

Again, if you need to change your answer, you may cross out the previous answer(s) and bubble in the new answer(s):

Select all that apply: Which are scientists?

- Stephen Hawking
- Albert Einstein
- Isaac Newton
- I don't know

For questions where you must fill in a blank, please make sure your final answer is fully included in the given space. You may cross out answers or parts of answers, but the final answer must still be within the given space.

Fill in the blank: What is the course number?

10-601

10-6301

Written Questions (31 points)

1 Latex Bonus Point (1 points)

1. (1 point) **Select one:** Did you use Latex for the entire written portion of this homework?

Yes

No

2 Warm-Up (11 points)

First, let's think a little bit about decision trees. The following dataset D consists of 8 examples, each with 3 attributes, (A, B, C) , and a label, Y .

A	B	C	Y
1	2	0	1
0	1	0	0
0	0	1	0
0	2	0	1
1	1	0	1
1	0	1	0
1	2	1	0
1	1	0	1

Use the data above to answer the following questions.

A few important notes:

- *All calculations should be done without rounding!* After you have finished all of your calculations, write your rounded solutions in the boxes below.
- Unless otherwise noted, numeric solutions should include 4 digits of precision (e.g. 0.1234).
- Note that, throughout this homework, we will use the convention that the leaves of the trees do not count as nodes, and as such are not included in calculations of depth and number of splits. (For example, a tree which classifies the data based on the value of a single attribute will have depth 1, and contain 1 split.)
- Note that the dataset contains duplicate rows; treat each of these as their own example, do not remove duplicate rows.

Note: Showing your work in these questions is optional, but it is recommended to help us understand where any misconceptions may occur. Only your numerical answer in the left box will be graded.

1. (1 point) What is the entropy of Y in bits, $H(Y)$? In this and subsequent questions, when we request the units in *bits*, this simply means that you need to use log base 2 in your calculations.¹ (Please include one number rounded to the fourth decimal place, e.g. 0.1234)

$H(Y)$	Work
1	$H(Y) = -(1/2)\log_2(1/2) - (1/2)\log_2(1/2)$

2. (1 point) What is the mutual information of Y and A in bits, $I(Y; A)$? (Please include one number rounded to the fourth decimal place, e.g. 0.1234)

$I(Y; A)$	Work
0.0488	$H(Y A=0) = -(1/3)\log_2(1/3) - (2/3)\log_2(2/3)$ $H(Y A=1) = -(2/5)\log_2(2/5) - (3/5)\log_2(3/5)$ $H(Y A) = (5/8) * H(Y A=1) + (3/8) * H(Y A=0) = 0.9512$ $I(Y; A) = H(Y) - H(Y A) = 0.04879$

3. (1 point) What is the mutual information of Y and B in bits, $I(Y; B)$? (Please include one number rounded to the fourth decimal place, e.g. 0.1234)

$I(Y; B)$
0.3113

¹If instead you used log base e , the units would be *nats*; log base 10 gives *bats*.

Work

$$\begin{aligned} H(Y|B=0) &= -(2/2)\log_2(2/2) = -1(0) = 0 \\ H(Y|B=1) &= -(1/3)\log_2(1/3) - (2/3)\log_2(2/3) \\ H(Y|B=2) &= -(1/3)\log_2(1/3) - (2/3)\log_2(2/3) \\ H(Y|B) &= 0 + (3/8)*H(Y|B=1) + (3/8)*H(Y|B=2) \\ I(Y;B) &= H(Y) - H(Y|B) = 0.311278 \end{aligned}$$

4. (1 point) What is the mutual information of Y and C in bits, $I(Y;C)$? (Please include one number rounded to the fourth decimal place, e.g. 0.1234)

$I(Y;C)$
0.5488

5. (1 point) Consider the dataset given above. Which attribute (A , B , or C) would a decision tree algorithm pick first to branch on, if its splitting criterion is mutual information?

Select one: :

- A
- B
- C

6. (1 point) Consider the dataset given above. After making the first split, which attribute would pick to branch on next, if the splitting criterion is mutual information? (Hint: Notice that this question correctly presupposes that there is *exactly one* second attribute.)


Select one:

- A
- B
- C

7. (1 point) If the same algorithm continues until the tree perfectly classifies the data, what would the depth of the tree be?

Depth
3

8. (4 points) Draw your completed Decision Tree. Label the non-leaf nodes with which attribute the tree will split on (e.g. B), the edges with the value of the attribute (e.g. 1 or 0), and the leaf nodes with the classification decision (e.g. $Y = 0$). You should include an image file below using the provided, commented out code in LaTeX, switching out *DecTree.png* to your file name as needed. The image may be hand-drawn.



3 Pseudocode (6 points)

1. In the programming assignment, you will need to implement three main tasks: training a decision tree on an arbitrary training set, predicting new values with a trained tree given an arbitrary input dataset, and evaluating your predictions against an arbitrary dataset's true labels. For this problem, we will focus on thinking through the algorithm for the *second* task.

Below, you will write pseudocode for the function `predict(node, example)`, which predicts the label of an `example` given a `node` of type `Node` representing the root of a trained tree. You must approach this problem recursively and use the `Node` class we have given to you.

```
class Node:  
    def __init__(self, attr, v):  
        self.attribute = attr  
        self.left = None  
        self.right = None  
        self.vote = v  
  
        # (a) the left and right children of a node are denoted as  
        #       node.left and node.right respectively, each is of  
        #       type Node  
        # (b) the attribute for a node is denoted as  
        #       node.attribute and has type str  
        # (c) if the node is a leaf, then node.vote of type str  
        #       holds the prediction from the majority vote; if node  
        #       is an internal node, then node.vote has value None  
        # (d) assume all attributes have values 0 and 1 only;  
        #       further assume that the left child corresponds to a  
        #       attribute value of 1, and the right child to a value  
        #       of 0  
  
    def predict(self, example):  
        # example is a dictionary which holds the attributes  
        # and the values of the attribute (ex. example["X"] =  
        # 0)
```

- (a) (3 points) Write the base case of `predict(node, example)`. Limit your answer to 10 lines.

Your Answer

```
if node.vote!=None:  
    return node.vote
```

- (b) (3 points) Write the recursive step of predict (node, example). Limit your answer to 10 lines.

Your Answer

```
value=example[node.attribute]
elif (value==1):
    return predict(node.left, example)
elif (value==0):
    return predict(node.right, example)
```


4 Empirical Questions (13 points)

The following questions should be completed as you work through the programming portion of this assignment.

1. (4 points) Train and test your decision tree on the politician dataset and the education dataset with four different values of max-depth, $\{0, 1, 2, 4\}$. Report your findings in the HW2 solutions template provided. A Decision Tree with max-depth 0 is simply a *majority vote classifier*; a Decision Tree with max-depth 1 is called a *decision stump*. If desired, you could even check that your answers for these two simple cases are correct using your favorite spreadsheet application (e.g. Excel, Google Sheets). (Please round each number to the fourth decimal place, e.g. 0.1234)

Dataset	Max-Depth	Train Error	Test Error
politician	0	0.4430	0.5060
	1	0.2013	0.2169
	2	0.1342	0.1566
	4	0.1074	0.2048
education	0	0.325	0.31
	1	0.195	0.23
	2	0.195	0.23
	4	0.13	0.155

2. (3 points) For the politicians dataset, create a computer-generated plot showing error on the y-axis against depth of the tree on the x-axis. Plot *both* training error and testing error, clearly labeling which is which. That is, for each possible value of max-depth (0, 1, 2, . . . , up to the number of attributes in the dataset), you should train a decision tree and report train/test error of the model's predictions. You should include an image file below using the provided, commented out code in LaTeX, switching out *politician.png* to your file name as needed.



3. (2 points) **Select one:** Suppose your research advisor asks you to run some model selection experiments and then report your results. You select the Decision Tree model's max-depth to be the one with lowest test error in metrics.txt and then report that model's test error as the performance of our classifier on held out test data. Is this a good experimental setup?

- Yes, because we are using the test set in order to choose the best model.
- No, because we should be using the *training* set to optimize the max-depth, not the test set.
- Yes, because we are not using the training set to tune hyperparameters.
- No, because we should be using a *validation* set to optimize the max-depth, not the test set.

4. (2 points) **Select one:** In this assignment, we used max-depth as our stopping criterion, and as a mechanism to prevent overfitting. Alternatively, we could stop splitting a node whenever the mutual information for the best attribute is lower than a threshold value. This threshold would be another hyperparameter. Theoretically, how would increasing this threshold value affect the number of nodes and depth of the learned trees?

- The higher this threshold value is, the *smaller* depth and the *fewer* nodes the decision tree contains.
- The higher this threshold value is, the *higher* depth and the *more* nodes the decision tree contains.
- The depth and number of nodes in the decision tree will not vary as the threshold value increases.

5. (2 points) Print (do not handwrite!) the decision tree which is produced by your algorithm for the politician data with max depth 3. Instructions on how to print the tree could be found in section 6.7.

Output

```
[83 democrat/66 republican]
| Superfund_right_to_sue = n: [55 democrat/2 republican]
| | Export_south_africa = n: [0 democrat/1 republican]
| | Export_south_africa = y: [55 democrat/1 republican]
| | | Immigration = n: [46 democrat/0 republican]
| | | Immigration = y: [9 democrat/1 republican]
| | Superfund_right_to_sue = y: [28 democrat/64 republican]
| | Aid_to_nicaraguan_contras = n: [13 democrat/58 republican]
| | | Export_south_africa = n: [0 democrat/20 republican]
| | | Export_south_africa = y: [13 democrat/38 republican]
| | Aid_to_nicaraguan_contras = y: [15 democrat/6 republican]
| | | Mx_missile = n: [12 democrat/0 republican]
| | | Mx_missile = y: [3 democrat/6 republican]
```

5 Collaboration Questions

After you have completed all other components of this assignment, report your answers to these questions regarding the collaboration policy. Details of the policy can be found [here](#).

1. Did you receive any help whatsoever from anyone in solving this assignment? If so, include full details.
2. Did you give any help whatsoever to anyone in solving this assignment? If so, include full details.
3. Did you find or come across code that implements any part of this assignment ? If so, include full details.

Your Answer

1)No 2)No 3)No

6 Programming: (70 points)

Your goal in this assignment is to implement a binary classifier, entirely from scratch—specifically a Decision Tree learner. In addition, we will ask you to run some end-to-end experiments on two tasks (predicting the party of a politician / predicting final grade for high school students) and report your results. You will write two programs: `inspection.py` (Section 6.2) and `decision_tree.py` (Section 6.3). The programs you write will be automatically graded using the Gradescope system. In this homework you have to choose Python as your programming language. Submitting code for more than one language may result in undefined behavior.

6.1 The Tasks and Datasets

Materials Download the zip file from the course website. The zip file will have a handout folder that contains all the data that you will need in order to complete this assignment.

Starter Code The handout will contain a preexisting `decision_tree.py` file that itself contains some starter code for the assignment. While we do not require that you use the starter code in your final submission, we *heavily* recommend building upon the structure layed out in the starter code.

Datasets The handout contains four datasets. Each one contains attributes and labels and is already split into training and testing data. The first line of each `.tsv` file contains the name of each attribute, and *the class is always the last column*.

1. **politician:** The first task is to predict whether a US politician is a member of the Democrat or Republican party, based on their past voting history. Attributes (aka. features) are short descriptions of bills that were voted on, such as *Aid_to_nicaraguan_contras* or *Duty_free_exports*. Values are given as ‘y’ for yes votes and ‘n’ for no votes. The training data is in `politicians_train.tsv`, and the test data in `politicians_test.tsv`.
2. **education:** The second task is to predict the final *grade* (A, not A) for high school students. The attributes (covariates, predictors) are student grades on 5 multiple choice assignments *M1* through *M5*, 4 programming assignments *P1* through *P4*, and the final exam *F*. The training data is in `education_train.tsv`, and the test data in `education_test.tsv`.
3. **small:** We also include `small_train.tsv` and `small_test.tsv`—a small, purely for demonstration version of the politicians dataset, with *only* attributes *Anti_satellite_test_ban* and *Export_south_africa*. For this small dataset, the handout tar file also contains the predictions from a reference implementation of a Decision Tree with max-depth 3 (see `small_3_train.labels`, `small_3_test.labels`, `small_3_metrics.txt`). You can check your own output against these to see if your implementation is correct.²

Note: For simplicity, all attributes are discretized into just two categories (i.e. each node will have at most two descendants). This applies to all the datasets in the handout, as well as the additional datasets on which we will evaluate your Decision Tree.

²Yes, you read that correctly: we are giving you the correct answers.

6.2 Program #1: Inspecting the Data [5pts]

Write a program `inspection.py` to calculate the label entropy at the root (i.e. the entropy of the labels before any splits) and the error rate (the percent of incorrectly classified instances) of classifying using a majority vote (picking the label with the most examples). You do not need to look at the values of any of the attributes to do these calculations, knowing the labels of each example is sufficient. **Entropy should be calculated in bits using log base 2.**

Command Line Arguments The autograder runs and evaluates the output from the files generated, using the following command:

```
$ python inspection.py <input> <output>
```

Your program should accept two command line arguments: an input file and an output file. It should read the `.tsv` input file (of the format described in Section 6.1), compute the quantities above, and write them to the output file so that it contains:

```
entropy: <entropy value>
error: <error value>
```

Example For example, suppose you wanted to inspect the file `small_train.tsv` and write out the results to `small_inspect.txt`. You would run the following command:

```
$ python inspection.py small_train.tsv small_inspect.txt
```

Afterwards, your output file `small_inspect.txt` should contain the following:

```
entropy: 0.996316519559
error: 0.464285714286
```

Our autograder will run your program on several input datasets to check that it correctly computes entropy and error, and will take minor differences due to rounding into account. You do not need to round your reported numbers! The Autograder will automatically incorporate the right tolerance for float comparisons.

For your own records, run your program on each of the datasets provided in the handout—this error rate for a *majority vote* classifier is a baseline over which we would (ideally) like to improve.

6.3 Program #2: Decision Tree Learner [65pts]

In `decision_tree.py`, implement a Decision Tree learner. This file should learn a decision tree with a specified maximum depth, print the decision tree in a specified format, predict the labels of the training and testing examples, and calculate training and testing errors.

Your implementation must satisfy the following requirements:

- Use mutual information to determine which attribute to split on.
- Be sure you're correctly weighting your calculation of mutual information. For a split on attribute X, $I(Y; X) = H(Y) - H(Y|X) = H(Y) - P(X=0)H(Y|X=0) - P(X=1)H(Y|X=1)$.
- As a stopping rule, only split on an attribute if the mutual information is > 0 .
- Do not grow the tree beyond a max-depth specified on the command line. For example, for a maximum depth of 3, split a node only if the mutual information is > 0 and the current level of the node is < 3 .
- Use a majority vote of the labels at each leaf to make classification decisions. If the vote is tied, choose the label that comes *last* in the lexicographical order (i.e. Republican should be chosen before Democrat)
- It is possible for different columns to have equal values for mutual information. In this case, you should split on the **first** column to break ties (e.g. if column 0 and column 4 have the same mutual information, use column 0).
- Do not hard-code any aspects of the datasets into your code. We may autograde your programs on hidden datasets that include different attributes and output labels.

Careful planning will help you to correctly and concisely implement your Decision Tree learner. Here are a few *hints* to get you started:

- Write helper functions to calculate entropy and mutual information.
- It is best to think of a Decision Tree as a collection of nodes, where nodes are either leaf nodes (where final decisions are made) or interior nodes (where we split on attributes). It is helpful to design a function to train a single node (i.e. a depth-0 tree), and then recursively call that function to create sub-trees.
- In the recursion, keep track of the depth of the current tree so you can stop growing the tree beyond the max-depth.
- Implement a function that takes a learned decision tree and data as inputs, and generates predicted labels. You can write a separate function to calculate the error of the predicted labels with respect to the given (ground-truth) labels.
- Be sure to correctly handle the case where the specified maximum depth is greater than the total number of attributes.
- Be sure to handle the case where max-depth is zero (i.e. a majority vote classifier).
- Look under the FAQ's on Piazza for more useful clarifications about the assignment.

6.4 Command Line Arguments

The autograder runs and evaluates the output from the files generated, using the following command:

```
$ python decision_tree.py [args...]
```

Where above [args...] is a placeholder for six command-line arguments: <train input> <test input> <max depth> <train out> <test out> <metrics out>. These arguments are described in detail below:

1. <train input>: path to the training input .tsv file (see Section 6.1)
2. <test input>: path to the test input .tsv file (see Section 6.1)
3. <max depth>: maximum depth to which the tree should be built
4. <train out>: path of output .labels file to which the predictions on the *training* data should be written (see Section 6.5)
5. <test out>: path of output .labels file to which the predictions on the *test* data should be written (see Section 6.5)
6. <metrics out>: path of the output .txt file to which metrics such as train and test error should be written (see Section 6.6)

As an example, the following command line would run your program on the politicians dataset and learn a tree with max-depth of two. The train predictions would be written to `pol_2_train.labels`, the test predictions to `pol_2_test.labels`, and the metrics to `pol_2_metrics.txt`.

```
$ python decision_tree.py politicians_train.tsv politicians_test.tsv \
  2 pol_2_train.labels pol_2_test.labels pol_2_metrics.txt
```

The following example would run the same learning setup except with max-depth three, and conveniently writing to analogously named output files, so you can compare the two runs.

```
$ python decision_tree.py politicians_train.tsv politicians_test.tsv \
  3 pol_3_train.labels pol_3_test.labels pol_3_metrics.txt
```

6.5 Output: Labels Files

Your program should write two output .labels files containing the predictions of your model on training data (<train out>) and test data (<test out>). Each should contain the predicted labels for each example printed on a new line. Use '\n' to create a new line.

Your labels should exactly match those of a reference decision tree implementation—this will be checked by the autograder by running your program and evaluating your output file against the reference solution.

Note: You should output your predicted labels using the same string identifiers as the original training data: e.g., for the politicians dataset you should output democrat/republican and for the education dataset you should output A/notA. The first few lines of an example output file is given below for the politician dataset:

```
democrat
democrat
democrat
republican
democrat
...
...
```

6.6 Output: Metrics File

Generate another file where you should report the training error and testing error. This file should be written to the path specified by the command line argument <metrics out>. Your reported numbers should be within 0.01 of the reference solution. You do not need to round your reported numbers! The Autograder will automatically incorporate the right tolerance for float comparisons. The file should be formatted as follows:

```
error(train): 0.0714
error(test): 0.1429
```

The values above correspond to the results from training a tree of depth 3 on `small_train.tsv` and testing on `small_test.tsv`. (There is one space between the colon and value)

6.7 Output: Printing the Tree

Finally, you should write a function to pretty-print your learned decision tree. **Your function should print your tree only after you are done generating the fully-trained tree.** Each row should correspond to a node in the tree. They should be printed using a *Pre-Order depth-first-search* traversal (but you may print left-to-right or right-to-left, i.e. your answer do not need to have exactly the same order as the reference below). Print the attribute of the node's parent and the attribute value corresponding to the node. Also include the sufficient statistics (i.e. count of positive / negative examples) for the data passed to that node. The row for the root should include *only* those sufficient statistics. A node at depth d , should be prefixed by d copies of the string '| '.

Below, we have provided the recommended format for printing the tree (example for python). You can print it directly to standard out rather than to a file. **This functionality of your program will not be autograded.**

```
$ python decision_tree.py small_train.tsv small_test.tsv 2 \
small_2_train.labels small_2_test.labels small_2_metrics.txt

[15 democrat/13 republican]
| Anti_satellite_test_ban = y: [13 democrat/1 republican]
| | Export_south_africa = y: [13 democrat/0 republican]
| | Export_south_africa = n: [0 democrat/1 republican]
| Anti_satellite_test_ban = n: [2 democrat/12 republican]
| | Export_south_africa = y: [2 democrat/7 republican]
| | Export_south_africa = n: [0 democrat/5 republican]
```

However, you should be careful that the tree might not be full. For example, after swapping the train/test files in the example above, you could end up with a tree like the following.

```
$ python decision_tree.py small_test.tsv small_train.tsv 2 \
swap_2_train.labels swap_2_test.labels swap_2_metrics.txt

[13 democrat/15 republican]
| Anti_satellite_test_ban = y: [9 democrat/0 republican]
| Anti_satellite_test_ban = n: [4 democrat/15 republican]
| | Export_south_africa = y: [4 democrat/10 republican]
| | Export_south_africa = n: [0 democrat/5 republican]
```

The following pretty-print shows the education dataset with max-depth 3. Use this example to check your

code before submitting your pretty-print of the politics dataset (asked in question 14 of the Empirical questions).

```
$ python decision_tree.py education_train.tsv education_test.tsv 3 \
edu_3_train.labels edu_3_test.labels edu_3_metrics.txt

[135 A/65 notA]
| F = A: [119 A/23 notA]
| | M4 = A: [56 A/2 notA]
| | | P1 = A: [41 A/0 notA]
| | | P1 = notA: [15 A/2 notA]
| | M4 = notA: [63 A/21 notA]
| | | M2 = A: [37 A/3 notA]
| | | M2 = notA: [26 A/18 notA]
| F = notA: [16 A/42 notA]
| | M2 = A: [13 A/15 notA]
| | | M4 = A: [6 A/1 notA]
| | | M4 = notA: [7 A/14 notA]
| | M2 = notA: [3 A/27 notA]
| | | M4 = A: [3 A/5 notA]
| | | M4 = notA: [0 A/22 notA]
```

The numbers in brackets give the number of positive and negative labels from the training data in that part of the tree.

At this point, you should be able to go back and answer questions 1-7 in the "Empirical Questions" section of this handout. Write your solutions in the template provided.

6.8 Submission Instructions

Programming Please ensure you have completed the following files for submission.

inspection.py
decision_tree.py

When submitting your solution, make sure to select and upload both files. Ensure the files have the exact same spelling and letter casing as above. You can either directly zip the two files (by selecting the two files and compressing them - do not compress the folder containing the files) or directly drag them to Gradescope for submission.

Written Questions Make sure you have completed all questions from Written component (including the collaboration policy questions) in the template provided. When you have done so, please submit your document in **pdf format** to the corresponding assignment slot on Gradescope.